

# EDA

August 11, 2025

#

**EDA**

```
[20]: # Import libraries
import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns
import tarfile
from sklearn.datasets import make_classification
from sklearn.preprocessing import (
    OneHotEncoder,
    OrdinalEncoder,
    LabelEncoder,
    StandardScaler,
    Normalizer,
)

# from sklearn.metrics import confusion_matrix,
# accuracy_score, plot_confusion_matrix, classification_report
from sklearn.metrics import (
    confusion_matrix,
    accuracy_score,
    ConfusionMatrixDisplay,
    classification_report,
)
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.linear_model import SGDClassifier

[21]: # Retrieve dataset
# Code developed with assistance from Generative AI

# check development environment
```

```

try:
    import google.colab
    IN_COLAB = True
except:
    IN_COLAB = False

# connect to drive if IN_COLAB = True
if IN_COLAB:
    from google.colab import drive
    drive.mount('/content/drive')

# set datapath
data_path = "../datasets/bank-additional-full.csv"

# download dataset if IN_COLAB = True
if IN_COLAB:
    print("Running in Colab... downloading dataset.")
    data_url = "https://raw.githubusercontent.com/junclemente/ads504-final-project/main/datasets/bank-additional-full.csv"
    if not os.path.exists("../datasets"):
        os.makedirs("../datasets")
    if not os.path.exists(data_path):
        !wget -q {data_url} -O {data_path}
else:
    print("Running locally... using local dataset.")

# Load the dataset
df = pd.read_csv(data_path, sep=";")

```

Running locally... using local dataset.

[22]: df.head()

```

[22]:   age      job  marital  education  default  housing  loan  contact  \
0   56  housemaid  married   basic.4y        no         no   no  telephone
1   57  services  married  high.school  unknown         no   no  telephone
2   37  services  married  high.school        no        yes   no  telephone
3   40   admin.  married   basic.6y        no         no   no  telephone
4   56  services  married  high.school        no         no   yes  telephone

   month  day_of_week  ...  campaign  pdays  previous  poutcome  emp.var.rate  \
0    may           mon  ...         1    999          0  nonexistent          1.1
1    may           mon  ...         1    999          0  nonexistent          1.1
2    may           mon  ...         1    999          0  nonexistent          1.1
3    may           mon  ...         1    999          0  nonexistent          1.1
4    may           mon  ...         1    999          0  nonexistent          1.1

```

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	93.994	-36.4	4.857	5191.0	no
1	93.994	-36.4	4.857	5191.0	no
2	93.994	-36.4	4.857	5191.0	no
3	93.994	-36.4	4.857	5191.0	no
4	93.994	-36.4	4.857	5191.0	no

[5 rows x 21 columns]

## 0.1 1. Basic Information

```
[23]: # Basic information
print("Shape of dataset:", df.shape)
print("\nData types:\n", df.dtypes)
```

Shape of dataset: (41188, 21)

Data types:

age	int64
job	object
marital	object
education	object
default	object
housing	object
loan	object
contact	object
month	object
day_of_week	object
duration	int64
campaign	int64
pdays	int64
previous	int64
poutcome	object
emp.var.rate	float64
cons.price.idx	float64
cons.conf.idx	float64
euribor3m	float64
nr.employed	float64
y	object

dtype: object

## 0.2 2. Missing Values and Unique Summary

```
[24]: # Missing values
print(df.isnull().sum())
```

age	0
job	0

marital	0
education	0
default	0
housing	0
loan	0
contact	0
month	0
day_of_week	0
duration	0
campaign	0
pdays	0
previous	0
poutcome	0
emp.var.rate	0
cons.price.idx	0
cons.conf.idx	0
euribor3m	0
nr.employed	0
y	0

dtype: int64

```
[25]: # Any duplicates?
      print(df.nunique())
```

age	78
job	12
marital	4
education	8
default	3
housing	3
loan	3
contact	2
month	10
day_of_week	5
duration	1544
campaign	42
pdays	27
previous	8
poutcome	3
emp.var.rate	10
cons.price.idx	26
cons.conf.idx	26
euribor3m	316
nr.employed	11
y	2

dtype: int64

### 0.3 3. Univariate Analysis (Numerical Features)

```
[26]: df.describe()
```

```
[26]:
```

	age	duration	campaign	pdays	previous \
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963
std	10.42125	259.279249	2.770014	186.910907	0.494901
min	17.00000	0.000000	1.000000	0.000000	0.000000
25%	32.00000	102.000000	1.000000	999.000000	0.000000
50%	38.00000	180.000000	2.000000	999.000000	0.000000
75%	47.00000	319.000000	3.000000	999.000000	0.000000
max	98.00000	4918.000000	56.000000	999.000000	7.000000

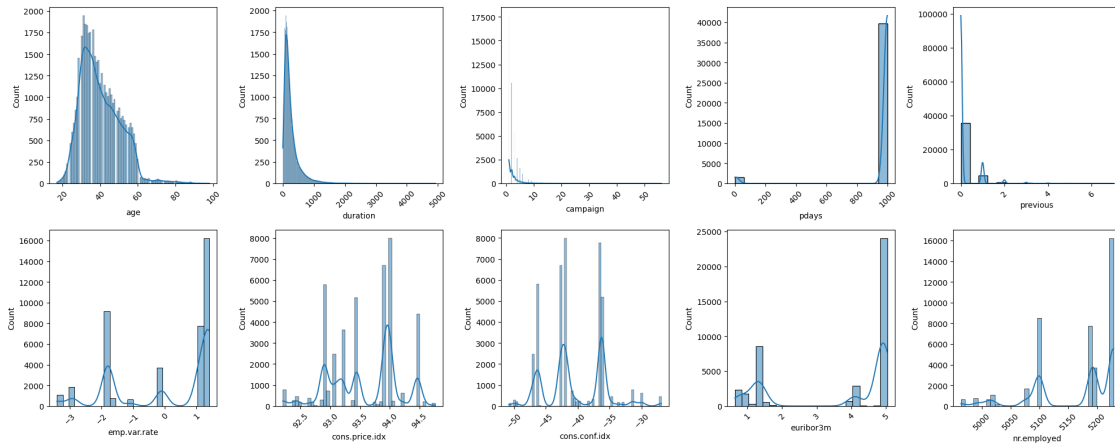
  

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	1.570960	0.578840	4.628198	1.734447	72.251528
min	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000

```
[27]: # Boxplot and histogram for a few important features
num_cols = df.select_dtypes(include="number").columns
n_cols = 5
n_rows = 2

fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, 8))
axes = axes.flatten()

for i, col in enumerate(num_cols[: n_cols * n_rows]):
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].tick_params(axis="x", rotation=45)
plt.tight_layout()
plt.show()
```



## 0.4 4. Univariate Analysis (Categorical Features)

```
[28]: # Summary statistics (categorical features)
print("\nSummary statistics (categorical):\n", df.describe(include="object"))
```

Summary statistics (categorical):

	job	marital	education	default	housing	loan	contact \
count	41188	41188	41188	41188	41188	41188	41188
unique	12	4	8	3	3	3	2
top	admin.	married	university.degree	no	yes	no	cellular
freq	10422	24928	12168	32588	21576	33950	26144

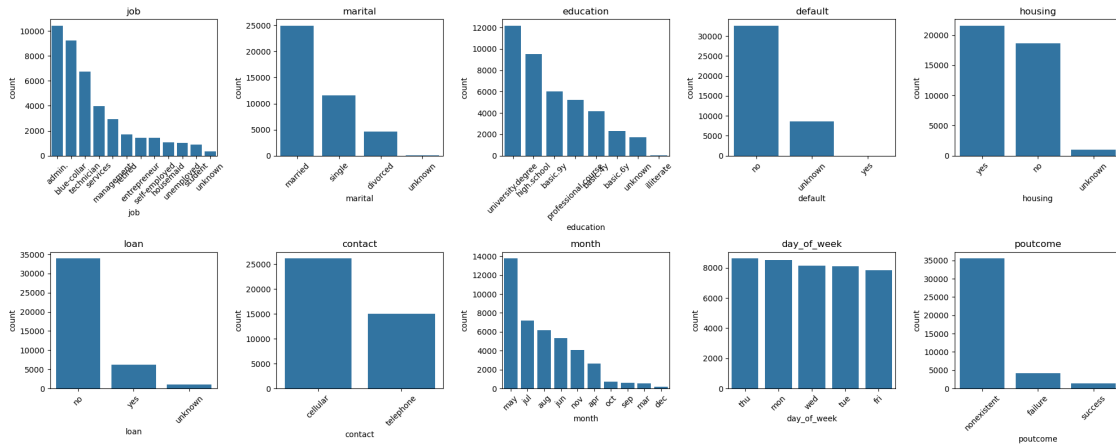
	month	day_of_week	poutcome	y
count	41188	41188	41188	41188
unique	10	5	3	2
top	may	thu	nonexistent	no
freq	13769	8623	35563	36548

```
[29]: cat_cols = df.select_dtypes(include="object").columns.drop("y")
n_cols = 5
n_rows = 2

fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, 8))
axes = axes.flatten()


for i, col in enumerate(cat_cols):
    sns.countplot(x=col, data=df, order=df[col].value_counts().index,
    ↪ax=axes[i])
    axes[i].set_title(f"{col}", fontsize=12)
    axes[i].tick_params(axis="x", rotation=45)
```

```
plt.tight_layout()
plt.show()
```



## 0.5 5. Target Variable (y) Distribution

```
[30]: # Plot target variable distribution
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x="y")
plt.title("Target Label Distribution")
plt.show()
```



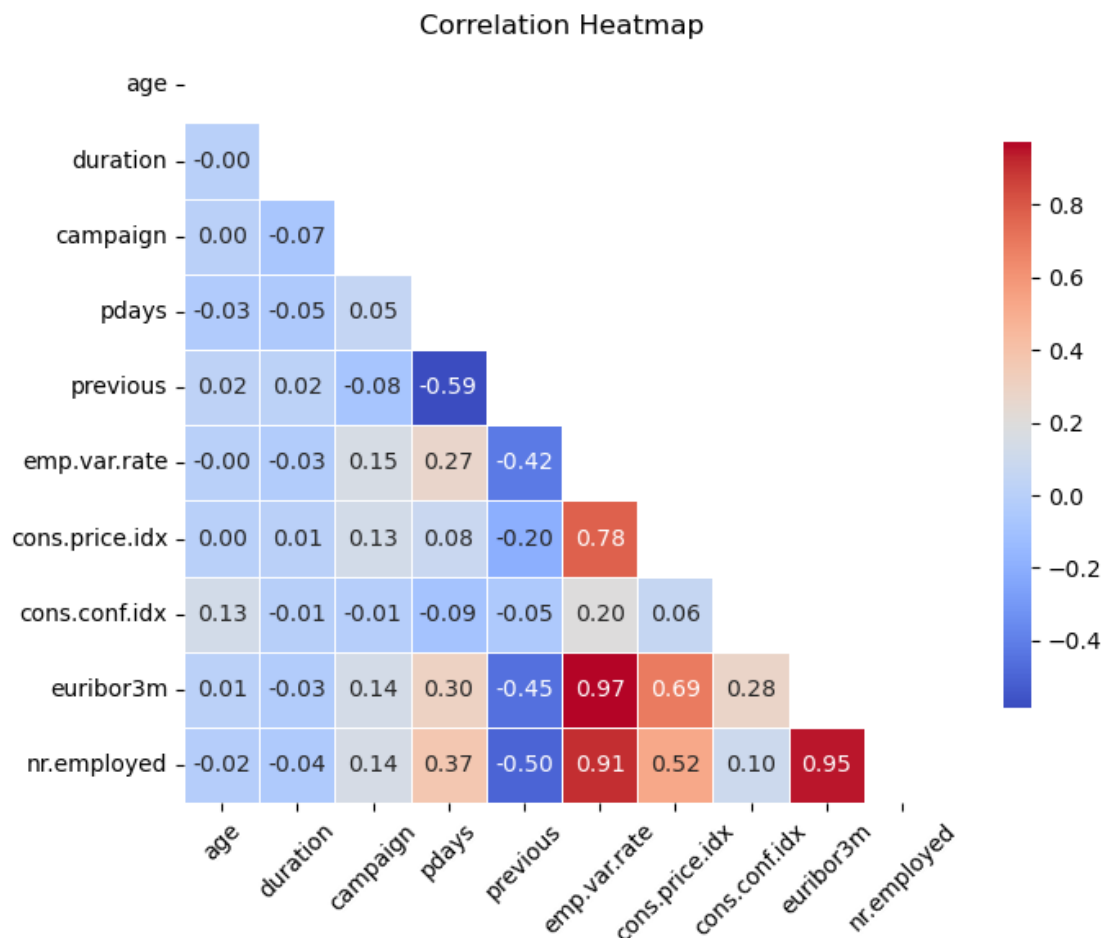
## 0.6 6. Correlation heatmap (only for numeric columns)

```
[31]: # Set matrix
corr = df.corr(numeric_only=True)
# Create a mask for the upper triangle plot
mask = np.triu(np.ones_like(corr, dtype=bool))

# set up the matplotlib figure
plt.figure(figsize=(10, 6))

sns.heatmap(
    corr,
    mask=mask,
    cmap="coolwarm",
    annot=True,
    fmt=".2f",
    square=True,
    linewidths=0.5,
    cbar_kws={"shrink": 0.75},
)
plt.xticks(rotation=45)
plt.title("Correlation Heatmap")
plt.show()
```





## 0.7 7. Correlation with Target label

```
[32]: # Encode target to numeric for correlation
df_temp = df.copy()
df_temp["y"] = df_temp["y"].map({"no": 0, "yes": 1})
print(df_temp.corr(numeric_only=True)["y"].sort_values(ascending=False))
```

```
y          1.000000
duration   0.405274
previous    0.230181
cons.conf.idx 0.054878
age         0.030399
campaign   -0.066357
cons.price.idx -0.136211
emp.var.rate -0.298334
euribor3m   -0.307771
pdays     -0.324914
nr.employed -0.354678
```

Name: y, dtype: float64

## 0.8 8. Bivariate Analysis: Categorical vs Target

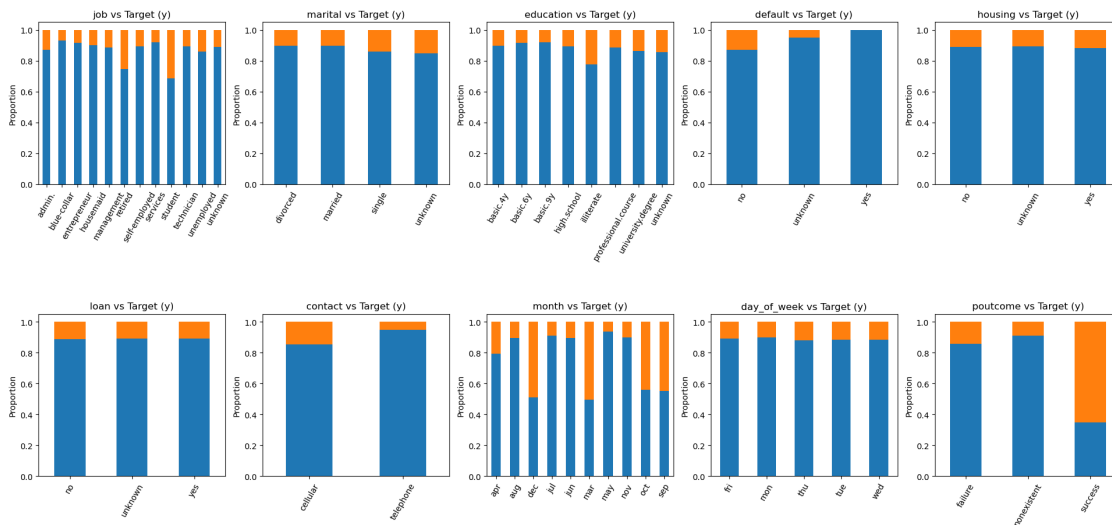
```
[33]: cat_cols = df.select_dtypes(include="object").columns.drop("y")
n_cols = 5
n_rows = 2

fig, axes = plt.subplots(n_rows, n_cols, figsize=(24, 10))
axes = axes.flatten()

for i, col in enumerate(cat_cols[: n_cols * n_rows]):
    ct = pd.crosstab(df[col], df["y"], normalize="index")
    ct.plot(kind="bar", stacked=True, ax=axes[i], legend=False)
    axes[i].set_title(f"{col} vs Target (y)", fontsize=12)
    axes[i].set_ylabel("Proportion")
    axes[i].tick_params(axis="x", rotation=60)
    axes[i].set_xlabel("")

fig.subplots_adjust(hspace=0.8, bottom=0.1)

plt.show()
```



### 0.8.1 Chi2 Contingency Table

```
[34]: from scipy.stats import chi2_contingency
from itertools import combinations
```

```

# get categorical columns
cat_cols = df.select_dtypes(include=["object", "category"]).columns.tolist()
# display(cat_cols)

# drop target column 'y'
cat_cols.remove("y")

# Cramer's V function
def cramer_v_score(x, y):
    table = pd.crosstab(x, y)
    chi2 = chi2_contingency(table)[0]
    n = table.sum().sum()
    return np.sqrt(chi2 / (n * (min(table.shape) - 1)))

results = []

for col in cat_cols:
    table = pd.crosstab(df[col], df["y"])

    # run chi2 test
    chi2, p, dof, expected = chi2_contingency(table)

    # cramer's v score
    cramers_score = cramer_v_score(df[col], df["y"])

    results.append(
        {
            "feature": col,
            "chi2_stat": chi2.round(4),
            "dof": dof,
            "p_value": p.round(4),
            "significant (p-value)": p < 0.05,
            "cramers_v": cramers_score.round(4),
        }
    )

results_df = pd.DataFrame(results)
results_df.sort_values(by="cramers_v", ascending=False, inplace=True)
display(results_df)


```

	feature	chi2_stat	dof	p_value	significant (p-value)	cramers_v
9	poutcome	4230.5238	2	0.0000	True	0.3205
7	month	3101.1494	9	0.0000	True	0.2744
0	job	961.2424	11	0.0000	True	0.1528
6	contact	862.3184	1	0.0000	True	0.1447
3	default	406.5775	2	0.0000	True	0.0994

2	education	193.1059	7	0.0000	True	0.0685
1	marital	122.6552	3	0.0000	True	0.0546
8	day_of_week	26.1449	4	0.0000	True	0.0252
4	housing	5.6845	2	0.0583	False	0.0117
5	loan	1.0940	2	0.5787	False	0.0052

Features with p-value < 0.05 are significantly associated with the target variable (y). These should be considered for further analysis or modeling.

```
[35]: plt.figure(figsize=(6, len(results_df) * 0.5))
sns.heatmap(
    results_df[["cramers_v"]].set_index(results_df["feature"]),
    annot=True,
    cmap="Blues",
    linewidths=0.5,
    cbar_kws={"label": "Cramér's V"},
)
plt.title("Cramér's V: Association Strength with Target (y)")
plt.xlabel("Strength")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



Features with Cramers  $V > 0.1$  are significantly associated with the target variable ( $y$ ). These should be considered for further analysis or modeling.