

Preprocess_Modeling

August 11, 2025

1 Pre-Processing of Bank Dataset

```
[1]: # Environment setup
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

from pathlib import Path

from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import FunctionTransformer, StandardScaler,
    ↳ OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import cross_validate
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score, roc_auc_score,
    confusion_matrix, ConfusionMatrixDisplay,
    roc_curve, RocCurveDisplay,
    precision_recall_curve, PrecisionRecallDisplay,
    average_precision_score
)

[2]: # install custom library
try:
    # check if jcds is installed
```

```

    from jcds import reports as jrep
except ImportError:
    print("Installing custom library: jcds")
    !pip install git+https://github.com/junclemente/jcgs.git

from jcds import reports as jrep
from jcds import eda as jeda

```

```

[3]: # Retrieve dataset
# Code developed with assistance from Generative AI

# check development environment
try:
    import google.colab
    IN_COLAB = True
except:
    IN_COLAB = False

# connect to drive if IN_COLAB = True
if IN_COLAB:
    from google.colab import drive
    drive.mount('/content/drive')

# set datapath
data_path = "../datasets/bank-additional-full.csv"

# download dataset if IN_COLAB = True
if IN_COLAB:
    print("Running in Colab... downloading dataset.")
    data_url = "https://raw.githubusercontent.com/junclemente/ads504-final_project/main/datasets/bank-additional-full.csv"
    if not os.path.exists("../datasets"):
        os.makedirs("../datasets")
    if not os.path.exists(data_path):
        !wget -q {data_url} -O {data_path}
else:
    print("Running locally... using local dataset.")

# Load the dataset
df = pd.read_csv(data_path, sep=";")
df.head()

```

Running locally... using local dataset.

```

[3]:   age      job  marital  education  default  housing  loan  contact  \
0    56  housemaid  married   basic.4y        no        no   no  telephone

```

1	57	services	married	high.school	unknown	no	no	telephone
2	37	services	married	high.school	no	yes	no	telephone
3	40	admin.	married	basic.6y	no	no	no	telephone
4	56	services	married	high.school	no	no	yes	telephone

	month	day_of_week	...	campaign	pdays	previous	poutcome	emp.var.rate	\
0	may	mon	...	1	999	0	nonexistent	1.1	
1	may	mon	...	1	999	0	nonexistent	1.1	
2	may	mon	...	1	999	0	nonexistent	1.1	
3	may	mon	...	1	999	0	nonexistent	1.1	
4	may	mon	...	1	999	0	nonexistent	1.1	

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	93.994	-36.4	4.857	5191.0	no
1	93.994	-36.4	4.857	5191.0	no
2	93.994	-36.4	4.857	5191.0	no
3	93.994	-36.4	4.857	5191.0	no
4	93.994	-36.4	4.857	5191.0	no

[5 rows x 21 columns]

1.1 Numerical Features

1.1.1 Cardinality Report

```
[4]: jrep.data_cardinality(df, show_columns=True)
```

CARDINALITY REPORT

Total columns analyzed: 21

[BINARY COLUMNS]

There are 2 binary columns.

* Columns: ['contact', 'y']

There are 0 binary with nan.

[CONSTANT/NEAR CONSTANT COLUMNS]

There are 0 constant columns.

There are 1 near-constant columns with $\geq 95\%$ of values being the same.

* Columns: ['pdays']

[LOW CARDINALITY CATEGORICAL COLUMNS]

* There are 10 low cardinality columns with ≤ 10 unique values.

Columns:

* marital: 4 unique values

* education: 8 unique values

* default: 3 unique values

* housing: 3 unique values

- * loan: 3 unique values
- * contact: 2 unique values
- * month: 10 unique values
- * day_of_week: 5 unique values
- * poutcome: 3 unique values
- * y: 2 unique values

[HIGH CARDINALITY CATEGORICAL COLUMNS]

- * There are 0 high cardinality variables with $\geq 90\%$ unique values.

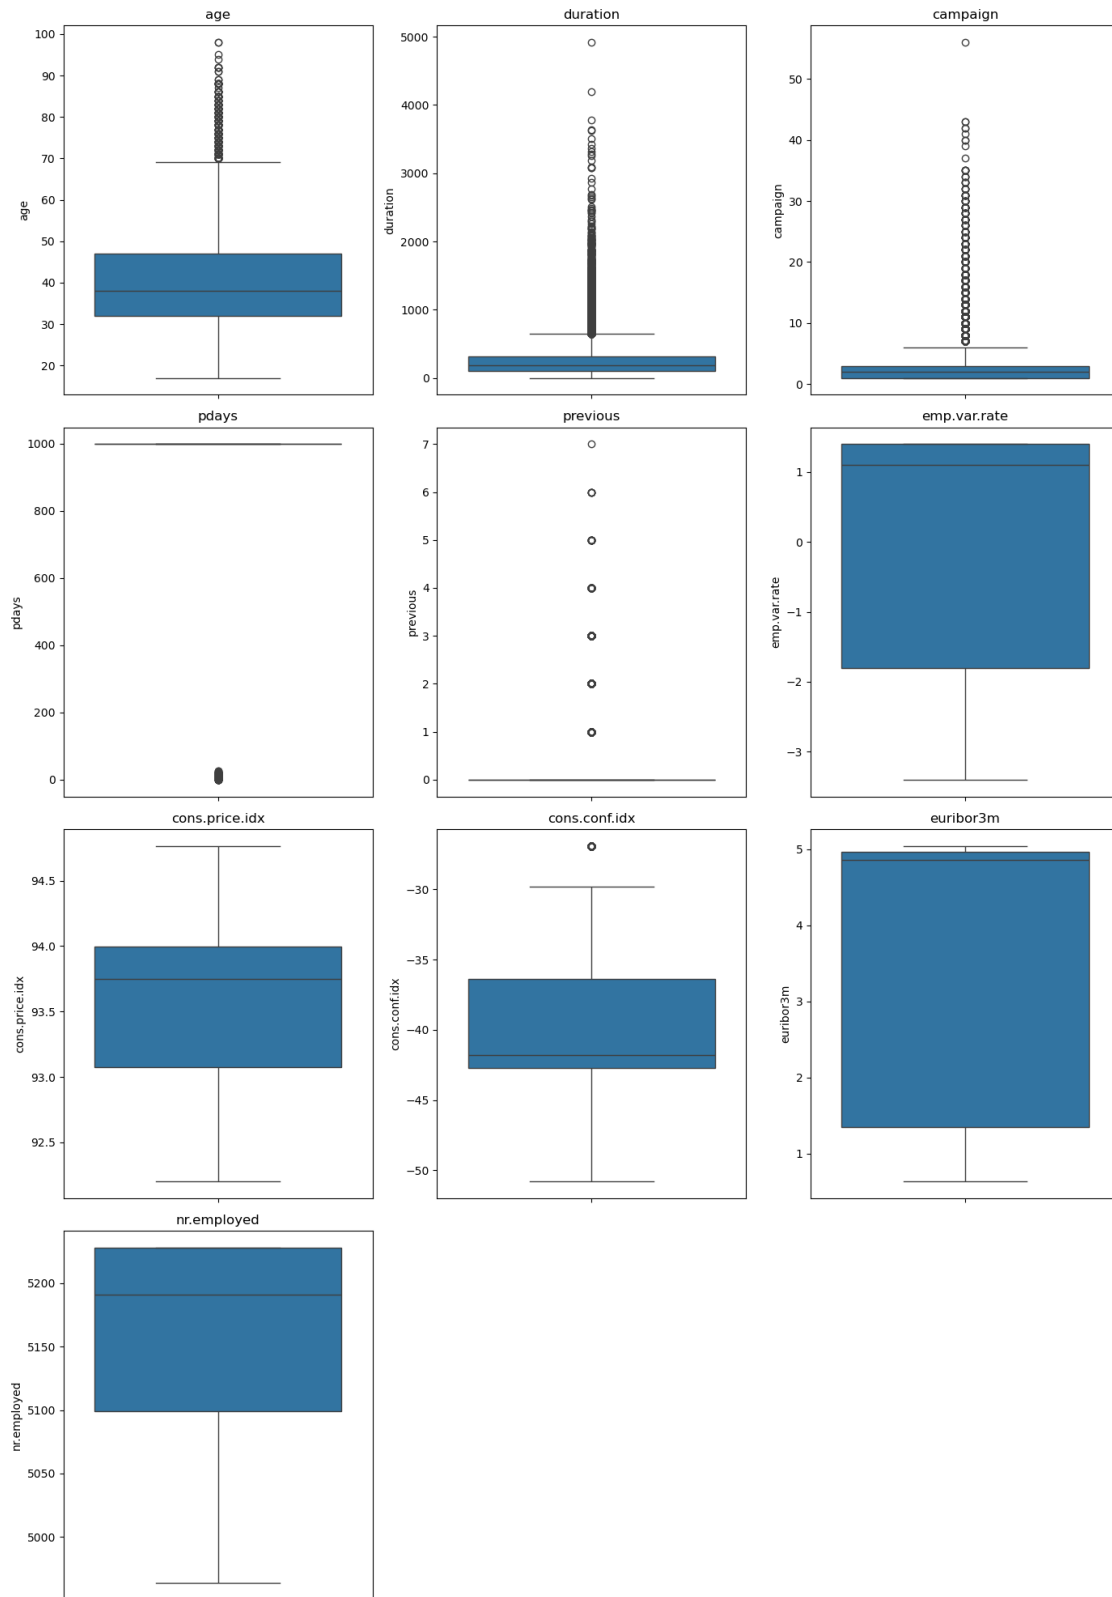
1.1.2 Outliers

```
[5]: # get total count of outliers
outliers = jeda.detect_outliers_iqr(df)

# create dataframe
outliers_df = pd.DataFrame(list(outliers.items()), columns=["feature", "count"])
display(outliers_df)
```

	feature	count
0	age	469
1	duration	2963
2	campaign	2406
3	pdays	1515
4	previous	5625
5	emp.var.rate	0
6	cons.price.idx	0
7	cons.conf.idx	447
8	euribor3m	0
9	nr.employed	0

```
[6]: jeda.plot_outlier_boxplots(df)
```



1.1.3 Observation

Six of the numerical variables have outliers.

1. **age outliers:** 70 - 100. These appear to be legitimate ages. No processing needed, maybe center/scale transform.
2. **duration:** based on data dictionary, this feature is highly correlated with target and should be removed.
3. **campaign:** There are 2406 detected outliers and showing a right skew. Transform to center/scale should be considered.
4. **pdays:** This could be converted to binary.
5. **previous:** Represents the number of times a customer was contacted prior to the campaign. This could be converted to binary: contacted vs not contacted.
6. **cons.conf.idx:** Total 447 outliers detected. Based on distribution from EDA, maybe this can be binned or turned to categorical.

Possible multicollinearity:

- euribor3m and emp.var.rate
- nr.employed and emp.var.rate
- cons.price.idx and emp.var.rate
- pdays and previous

1.2 Categorical Features

1.2.1 Observations

- marital: does not show much variability; may not be a good predictor
- housing: does not show much variability; not a good predictor unless **unknown** could have information
- loan: same as housing, does not show much variability. Also has a **unknown**
- day_of_week: does not show much variability; may not be a good predictor

Based on chi_2 stat, **housing** and **loan** are not statistically significant with p-value > 0.05.

With the Cramer's V statistic, **marital** and **day_of_week** do not show strong association with the target, 0.5, and 0.3 respectively. These may not be good predictors and could be candidates for removal from the final dataset for modeling.

1.3 Variables to Drop

Due to multicollinearity with **emp.var.rate**, the following numerical features will be dropped:

- euribor3m
- nr.employed
- cons.price.idx

*Out of these four numerical features, without someone with domain knowledge, generative AI was asked which feature would be best to keep. ChatGPT recommended keeping **emp.var.rate** due to "strong economic indicator, lower redundancy, and most interpretable."*

pdays and previous also show colinearity. Consider converting pdays to binary and dropping previous as it is not a strong predictor of the target.

Due to low statistical significance and low strength of association, the following categorical features will be dropped:

- housing
- loan

2 Transformation Pipelines

```
[7]: jrep.data_info(df, show_columns=True)
```

SHAPE:

There are 41188 rows and 21 columns (30.26 MB).

DUPLICATES:

There are 12 duplicated rows.

COLUMNS/VARIABLES:

Column dtype Summary:

- * object: 11
- * int: 5
- * float: 5

There are 10 numerical (int/float/bool) variables.

* Columns: ['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']

There are 11 categorical (nominal/ordinal) variables.

* Columns: ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcome', 'y']

DATETIME COLUMNS:

There are 0 datetime variables and 0 possible datetime variables.

OTHER COLUMN/VARIABLE INFO:

ID Like Columns (threshold = 95.0%): 0

Columns with mixed datatypes: 0

- * Columns: []

```
[8]: # get feature list by type
cat_var = jeda.show_catvar(df)
con_var = jeda.show_convar(df)
target = "y"

# remove target from categorical list
cat_var = [col for col in cat_var if col != target]
```

```
print(f"Categorical variables:", cat_var)
print(f"Continuous variables:", con_var)
```

Categorical variables: ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'poutcome']
 Continuous variables: ['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']

```
[9]: cols_to_drop = [
    "housing",
    "loan",
    "previous",
    "pdays", # after converting to binary
    "euribor3m",
    "nr.employed",
    "cons.price.idx",
]

con_var = [c for c in con_var if c not in cols_to_drop]
cat_var = [c for c in cat_var if c not in cols_to_drop]
```

2.1 Train / Test Split

```
[10]: # Encode target variable to 0, 1 prior to train_test_split
df[target] = df[target].replace({"no": 0, "yes": 1}).astype(int)
# show target unique values
jeda.list_unique_values(df, column=target)
```

```
>>> EXECUTING DataFrame["y"].unique().tolist()
```

```
Unique values in 'y':
[0, 1]
```

```
/tmp/ipykernel_102333/570709414.py:2: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df[target] = df[target].replace({"no": 0, "yes": 1}).astype(int)
```

```
[11]: RANDOM_STATE = 42
X = df.drop(columns=target)
y = df[target]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ stratify = y, random_state = RANDOM_STATE)
```


2.2 Define Pipeline Tranformeers

```
[12]: # columns to drop
def drop_cols(X):
    X = X.copy()
    return X.drop(columns=cols_to_drop)

drop_columns = FunctionTransformer(drop_cols, validate=False)

# convert pdays to binary
def pdays_to_binary(X):
    X = X.copy()
    X["pdays_binary"] = (X["pdays"] != 999).astype(int)
    return X

convert_to_binary = FunctionTransformer(pdays_to_binary, validate=False)

# scaler and encoder
preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), con_var),
        ("cat", OneHotEncoder(handle_unknown="ignore"), cat_var),
    ]
)
```

2.3 Modeling on the Training Set

```
[13]: # define function for modeling generator
def make_pipeline(clf):
    return ImbPipeline(steps=[
        ("bin", convert_to_binary),
        ("drop", drop_columns),
        ("encoder", preprocessor),
        ("smote", SMOTE(random_state=RANDOM_STATE)),
        ("classifier", clf)
    ])

]
```

```
[14]: # initialize the models (classifiers)
models = {
    "Perceptron": Perceptron(random_state=RANDOM_STATE),
    "Logistic Regression": LogisticRegression(random_state=RANDOM_STATE),
    "K-Nearest Neighbors": KNeighborsClassifier(),
    "Support Vector Machine": SVC(random_state=RANDOM_STATE, probability=True),
    "Neural Network": MLPClassifier(random_state=RANDOM_STATE, max_iter=1000),
    "XGBoost": XGBClassifier(random_state=RANDOM_STATE,
        use_label_encoder=False, eval_metric='logloss'),
}
```

```

    "AdaBoost": AdaBoostClassifier(random_state=RANDOM_STATE),
    "Random Forest": RandomForestClassifier(random_state=RANDOM_STATE)
}

```

```

[15]: # set scoring metrics
scoring = ["accuracy", "precision", "recall", "f1", "roc_auc"]

```

```

[ ]: # cross-validate each model
# initialize results dictionary
avg_train_results = {}
full_train_results = {}
for name, clf in models.items():
    pipe = make_pipeline(clf)
    cv_results = cross_validate(pipe, X_train, y_train, cv=5, scoring=scoring,
    ↪n_jobs=4, return_train_score=True)
    full_train_results[name] = {
        "fit_time": cv_results["fit_time"],
        "score_time": cv_results["score_time"],
        "train_accuracy": cv_results["train_accuracy"],
        "train_precision": cv_results["train_precision"],
        "train_recall": cv_results["train_recall"],
        "train_f1": cv_results["train_f1"],
        "train_roc_auc": cv_results["train_roc_auc"]}
    avg_train_results[name] = {metric: np.mean(cv_results[f"train_{metric}"])}
    ↪for metric in scoring}
    avg_train_results[name].update({f"std_{metric}": np.
    ↪std(cv_results[f"train_{metric}"])} for metric in scoring})

```

```

/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```

```

    bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```

```

    bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:07] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:08] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-
packages/xgboost/training.py:183: UserWarning: [23:42:08] WARNING:
/home/conda/feedstock_root/build_artifacts/xgboost-
split_1748293041487/work/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
[18]: # create dataframe of average results
avg_train_results_df = pd.DataFrame(avg_train_results).T
# round results to 3 decimal places
avg_train_results_df = avg_train_results_df.round(3)
# display average results
display(avg_train_results_df)
```

	accuracy	precision	recall	f1	roc_auc	\
Perceptron	0.815	0.356	0.757	0.482	0.872	
Logistic Regression	0.861	0.441	0.868	0.585	0.933	
K-Nearest Neighbors	0.901	0.533	0.998	0.695	0.997	

Support Vector Machine	0.897	0.525	0.934	0.672	0.962
Neural Network	0.975	0.849	0.941	0.893	0.994
XGBoost	0.937	0.729	0.696	0.712	0.967
AdaBoost	0.878	0.476	0.807	0.599	0.929
Random Forest	1.000	1.000	1.000	1.000	1.000

	std_accuracy	std_precision	std_recall	std_f1	\
Perceptron	0.024	0.037	0.061	0.032	
Logistic Regression	0.001	0.002	0.002	0.002	
K-Nearest Neighbors	0.001	0.003	0.001	0.003	
Support Vector Machine	0.001	0.004	0.001	0.003	
Neural Network	0.004	0.022	0.006	0.014	
XGBoost	0.001	0.006	0.008	0.006	
AdaBoost	0.002	0.006	0.011	0.005	
Random Forest	0.000	0.000	0.000	0.000	

	std_roc_auc
Perceptron	0.020
Logistic Regression	0.001
K-Nearest Neighbors	0.000
Support Vector Machine	0.001
Neural Network	0.001
XGBoost	0.001
AdaBoost	0.002
Random Forest	0.000

2.4 Out of Sample Validation Testing

```
[19]: # initialize test metrics and predictions
metrics = []
preds = {} # preds[name] = (y_test, y_pred, y_proba)

for name, clf in models.items():
    pipe = make_pipeline(clf)
    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)
    # try proba or decision_function
    if hasattr(pipe, 'predict_proba'):
        y_score = pipe.predict_proba(X_test)[:,-1]
    else:
        y_score = pipe.decision_function(X_test)

    metrics.append({
        'model': name,
        'accuracy': accuracy_score(y_test, y_pred),
        'precision': precision_score(y_test, y_pred),
        'recall': recall_score(y_test, y_pred),
```

```

        'f1':          f1_score(y_test,      y_pred),
        'roc_auc':     roc_auc_score(y_test, y_score)
    })
    preds[name] = (y_test, y_pred, y_score)

test_df = pd.DataFrame(metrics).set_index('model').round(3)
display(test_df)

```

/home/junc/miniconda3/envs/ads504_project/lib/python3.10/site-packages/xgboost/training.py:183: UserWarning: [23:49:12] WARNING: /home/conda/feedstock_root/build_artifacts/xgboost-split_1748293041487/work/src/learner.cc:738: Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

	accuracy	precision	recall	f1	roc_auc
model					
Perceptron	0.849	0.404	0.721	0.518	0.892
Logistic Regression	0.866	0.452	0.899	0.602	0.940
K-Nearest Neighbors	0.862	0.439	0.820	0.572	0.889
Support Vector Machine	0.881	0.483	0.871	0.622	0.943
Neural Network	0.892	0.519	0.558	0.538	0.911
XGBoost	0.916	0.631	0.616	0.624	0.949
AdaBoost	0.884	0.491	0.815	0.613	0.934
Random Forest	0.913	0.607	0.652	0.629	0.944

2.4.1 Modeling Performance Visualizations

Confusion Matrices

```

[20]: # preds[name] == (y_true, y_pred, y_score)
model_names = list(preds.keys())
n_models    = len(model_names)

# choose grid size
n_cols = 4
n_rows = int(np.ceil(n_models / n_cols))

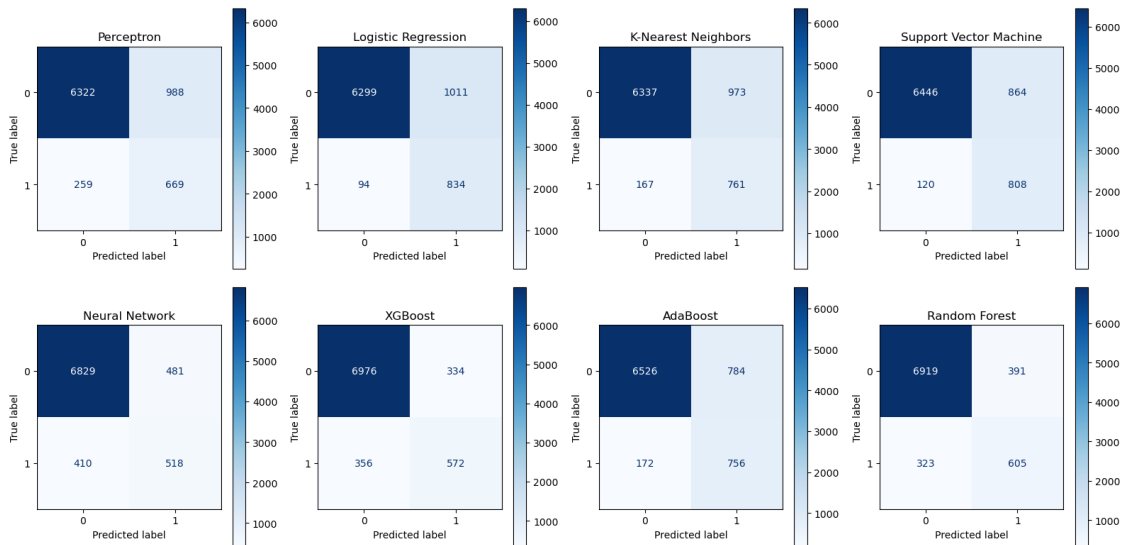
fig, axes = plt.subplots(n_rows, n_cols, figsize=(n_cols*4, n_rows*4))
axes = axes.flatten()

for ax, name in zip(axes, model_names):
    y_true, y_pred, _ = preds[name]
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(cm, display_labels=[0,1])
    disp.plot(ax=ax, cmap="Blues", values_format="d")
    ax.set_title(name)

```

```
# hide any unused subplots
for ax in axes[n_models:]:
    ax.axis("off")

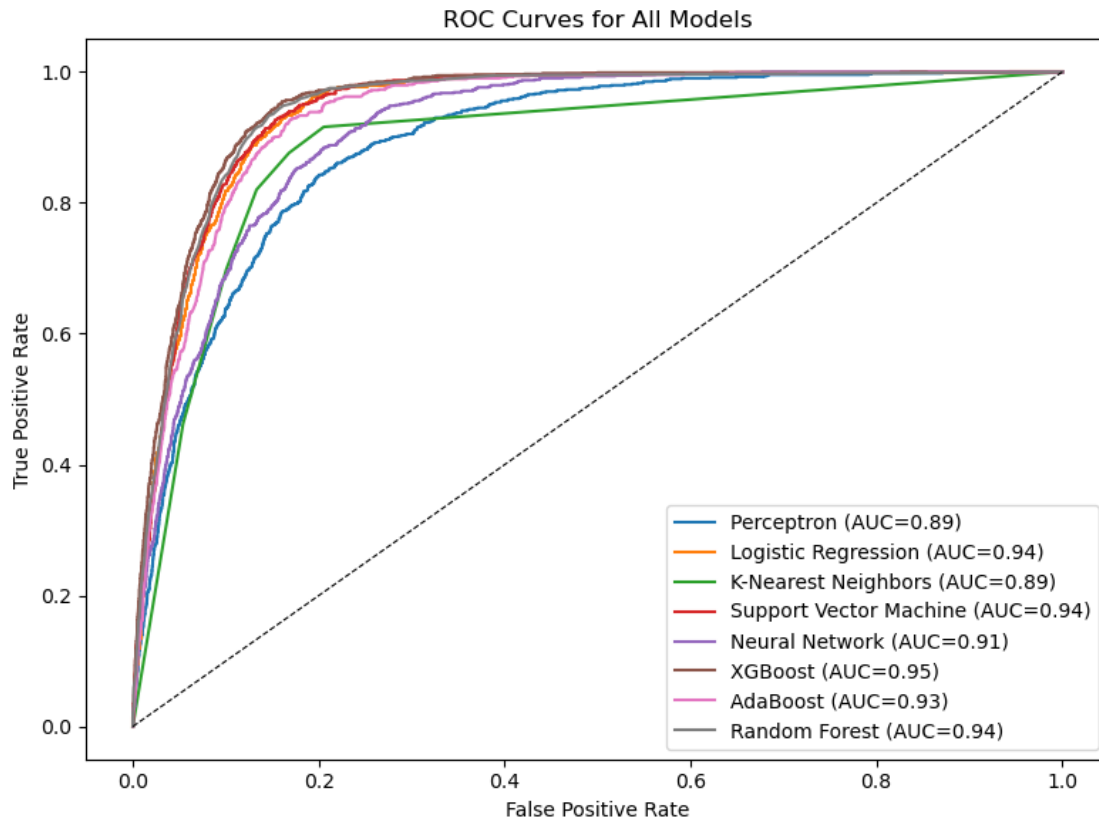
plt.tight_layout()
plt.show()
```



ROC-AUC Chart

```
[21]: fig, ax = plt.subplots(figsize=(8,6))
for name, (y_t, _, y_score) in preds.items():
    fpr, tpr, _ = roc_curve(y_t, y_score)
    ax.plot(fpr, tpr, label=f"{name} (AUC={roc_auc_score(y_t, y_score):.2f})")

ax.plot([0,1], [0,1], 'k--', linewidth=0.8)
ax.set_xlabel("False Positive Rate")
ax.set_ylabel("True Positive Rate")
ax.set_title("ROC Curves for All Models")
ax.legend(loc='lower right')
plt.tight_layout()
plt.show()
```



Precision Recall Curves

```
[22]: fig, ax = plt.subplots(figsize=(8,6))
for name, (y_t, _, y_score) in preds.items():
    precision, recall, _ = precision_recall_curve(y_t, y_score)
    ap = average_precision_score(y_t, y_score)
    ax.plot(recall, precision, label=f"{name} (AP={ap:.2f})")

ax.set_xlabel("Recall")
ax.set_ylabel("Precision")
ax.set_title("Precision-Recall Curves for All Models")
ax.legend(loc='upper right')
plt.tight_layout()
plt.show()
```

